Project Presentation

Moving Object Detector

CS677 Parallel Computing for Many Core Processors

Adrien HUET

Problem Description

- Implement a program capable of delimiting an moving object in a video input

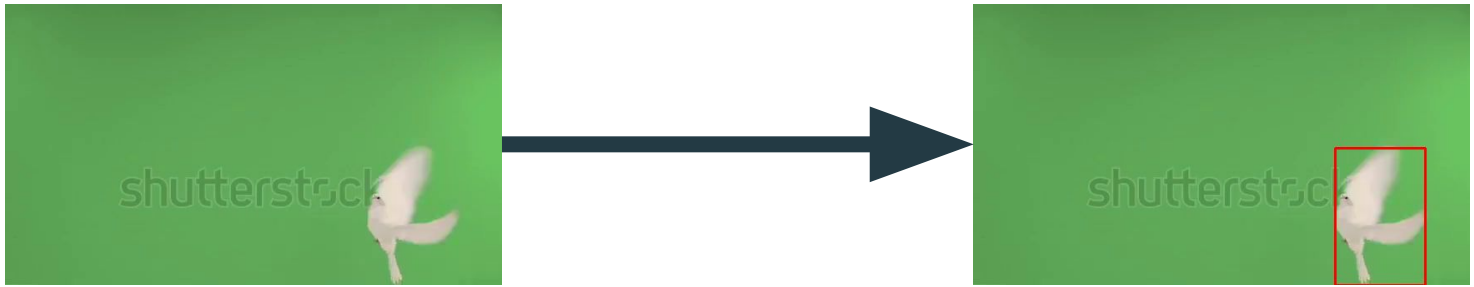

Algorithm

(1) Reserve the first frame *g0* of the video as a background reference for further frames.
(2) Convert both the background *g0* and the frame to process *g1* into grayscale.
(3) Smooth the two images *g0* and *g1* (e.g. with a Gaussian filter).
(4) Compute the difference $d = |g0 - g1|$ between the two images.
(5) Perform morphological closing/opening to remove non-meaningful artifacts.
(6) Threshold the image *d* and keep only the connected components (label propagation).
(7) Compute and output the bounding boxes.
(8) (Render new video with boxes drawn onto each frame)

Algorithm

Background

Frame

Algorithm

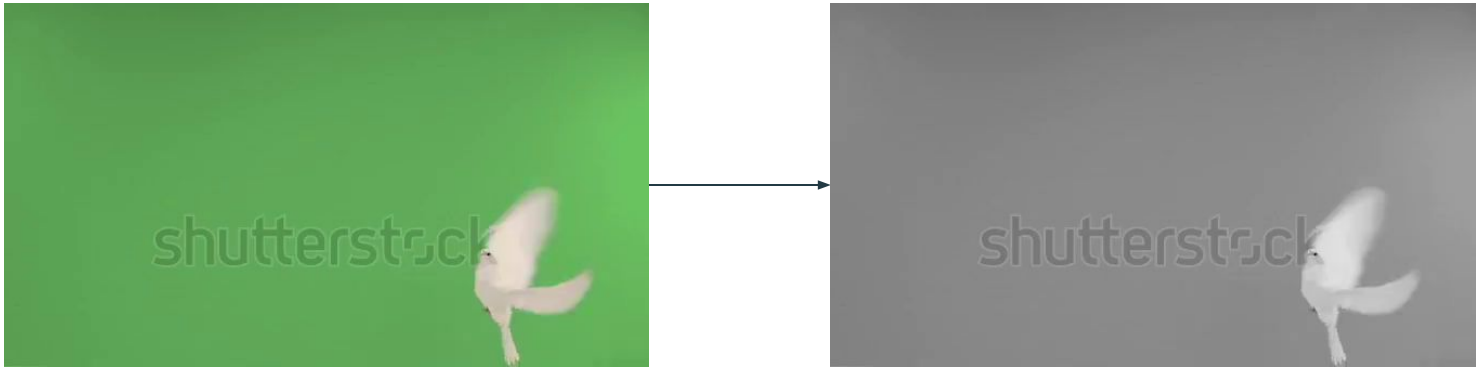

Frame

Grayscale

Algorithm

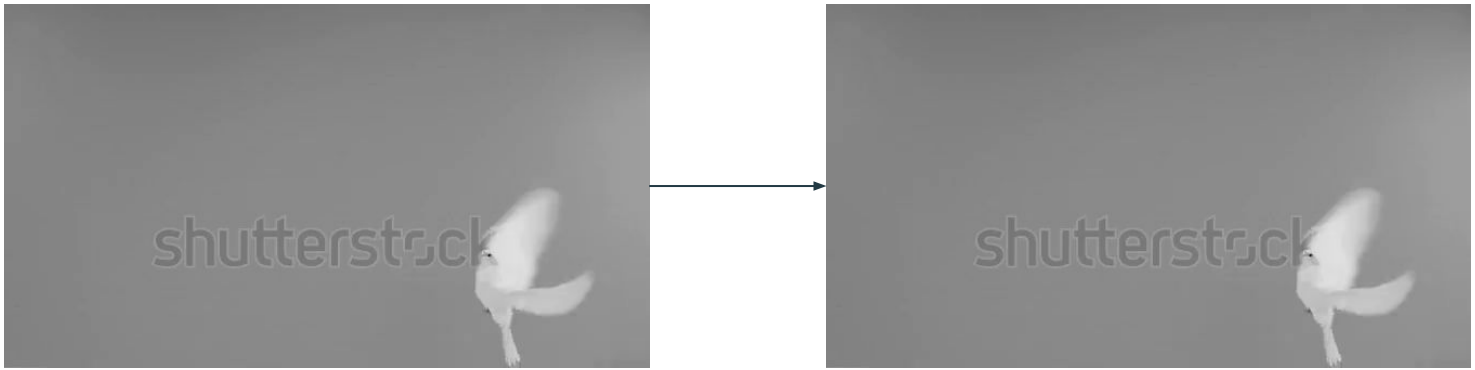

Grayscale

Blur

Algorithm

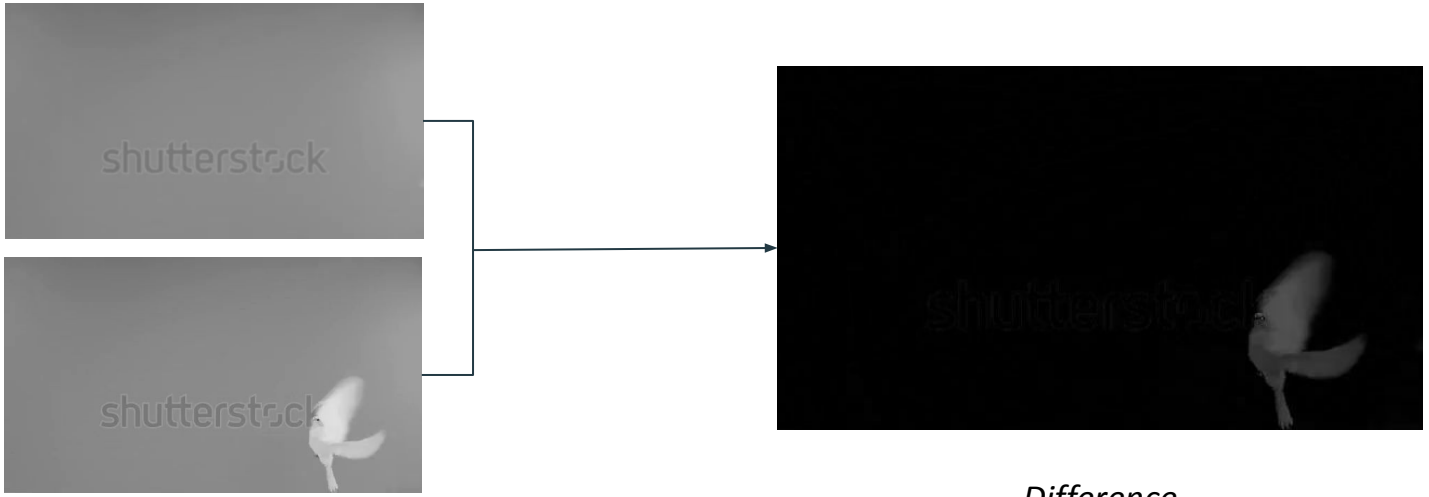

Difference

Algorithm

Difference

Morphological opening

Algorithm

Morphological opening

Threshold

Algorithm

Threshold

Symbolic CC
Labelled Image
+
Bounding Box

Connected Components Labelling - Union Find

"Optimizing GPU-Based Connected Components Labeling Algorithms", Stefano Allegretti, Federico Bolelli *et al.*, 2018, IPAS
"Optimized Block-Based Algorithms to Label Connected Components on GPUs", Stefano Allegretti, Federico Bolelli *et al.,* 2020, IEEE

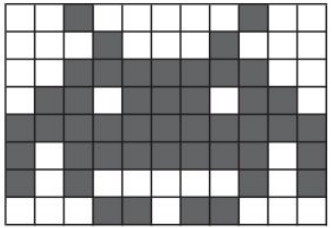

(a) Binary Input

1	2	0	4	5	6	7	8	0	10	11
12	13	14	0	16	17	18	0	20	21	22
23	24	0	0	0	0	0	0	0	32	33
34	0	0	37	0	0	0	41	0	0	44
0	0	0	0	0	0	0	0	0	0	0
0	57	0	0	0	0	0	0	0	65	0
0	68	0	70	71	72	73	74	0	76	0
78	79	80	0	0	83	0	0	86	87	88

(b) Output Initialization

1	1	0	4	4	4	4	4	0	10	10
1	1	14	0	16	17	18	0	20	21	22
23	24	0	0	0	0	0	0	0	32	33
34	0	0	37	0	0	0	41	0	0	44
0	0	0	0	0	0	0	0	0	0	0
0	57	0	0	0	0	0	0	0	65	0
0	68	0	70	71	72	73	74	0	76	0
78	79	80	0	0	83	0	0	86	87	88

(c) Provisional Result

1	1	0	1	1	1	1	1	0	1	1
1	1	1	0	1	1	1	0	1	1	1
1	1	0	0	0	0	0	0	0	1	1
1	0	0	37	0	0	0	41	0	0	1
0	0	0	0	0	0	0	0	0	0	0
0	57	0	0	0	0	0	0	0	57	0
0	57	0	57	57	57	57	57	0	57	0
57	57	57	0	0	57	0	0	57	57	57

(d) Output Labels

Suitability for GPU acceleration

- All steps can be executed on the GPU
- Morphological operations, blurring and finding connected components are the most expensive
- Low data footprint: predicted 1 load for background and 1 for each frame

GPU Implementation

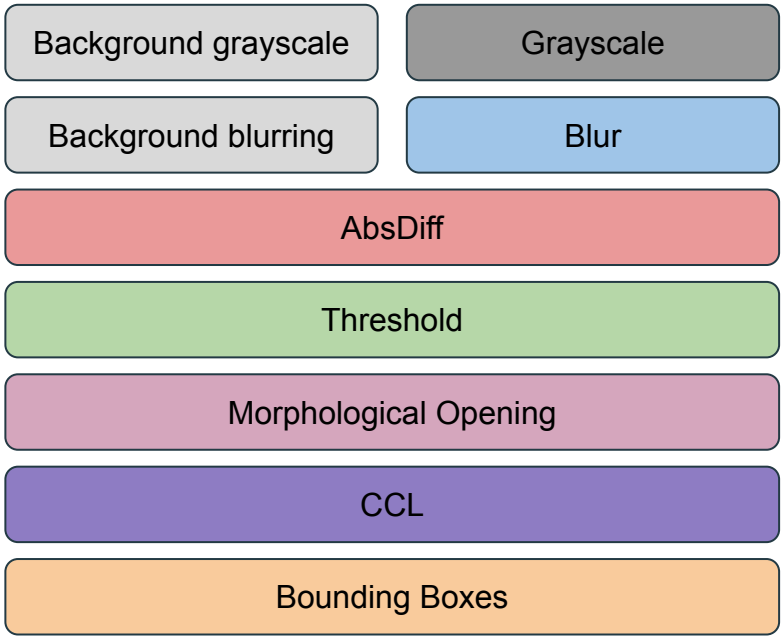

GPU Implementation

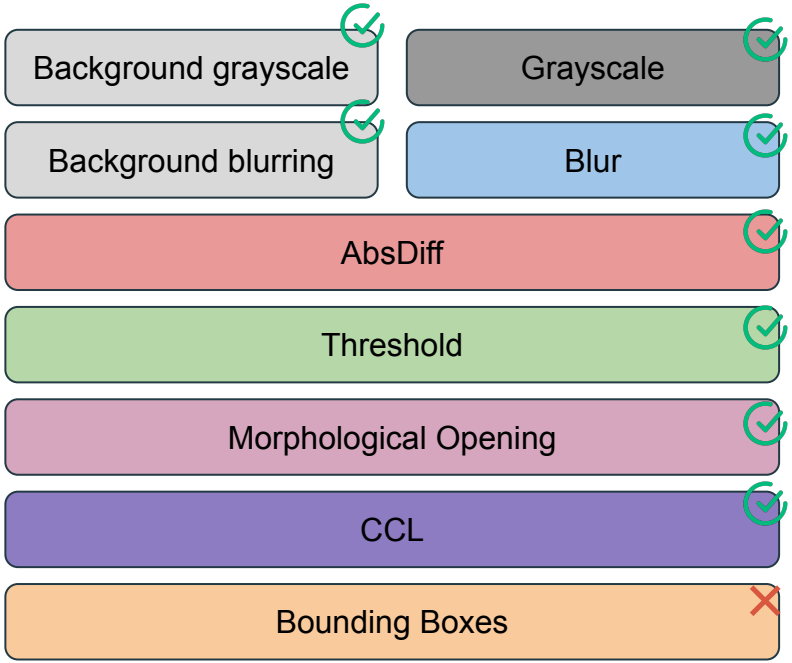

GPU V1.0

Very naive implementation

Allocation and deallocation for each frame

Background and kernels are processed for each frame

```
--------------------------------------------------------------------------------
OPENCV Bench (v1):
FPS: 450.4

STEP                    |      FRAME_AVG           TOTAL        EXEC_TIME

grayscale               |        215.75µs        26753µs            9.53%
blur                    |      303.1129µs        37586µs            13.4%
diff                    |      15.59677µs         1934µs           0.689%
threshold               |      11.12903µs         1380µs           0.492%
morph                   |      1127.669µs       139831µs            49.8%
connectedComps          |      190.7661µs        23655µs            8.43%
bboxes                  |      109.2258µs        13544µs            4.82%

Start to finish: 281ms
--------------------------------------------------------------------------------
CPU Bench (v1):
FPS: 2.61

STEP                    |      FRAME_AVG           TOTAL        EXEC_TIME

grayscale               |     0.9919355ms          123ms           0.259%
blur                    |      239.0323ms        29640ms          0.0623%
  - getGaussianMatrix   |      11.29032µs         1400µs           2.94%
diff                    |      240.1371µs        29777µs          0.0626%
threshold               |       260.629µs        32318µs          0.0679%
morph                   |      110.0645ms        13648ms            28.7%
  - getCircleKernel     |     0.7016129µs           87µs        0.000183%
connectedComps          |      14.17742ms         1758ms            3.7%
bboxes                  |      15.82258ms         1962ms           4.12%

Start to finish: 47.6s
--------------------------------------------------------------------------------
GPU Bench (v1):
blockDim: 32x32
gridDim: 19x11
FPS: 101.7

STEP                    |      FRAME_AVG           TOTAL        EXEC_TIME

grayscaleGPU            |    0.05608258ms        6.95424ms           0.569%
blurGPU                 |       3.045465ms      377.6376ms            30.9%
  - getGaussianMatrix   |    0.001887226ms       0.234016ms          0.0191%
diffGPU                 |    0.02317884ms        2.874176ms           0.235%
thresholdGPU            |    0.02207355ms         2.73712ms           0.224%
morph                   |       4.239244ms      525.6662ms             43%
  - dilateGPU           |       2.104713ms      260.9844ms            21.4%
  - erodeGPU            |       2.111592ms      261.8374ms            21.4%
  - getCircleKernel     |      0.001336ms       0.165664ms          0.0136%
connectedComps          |      0.5183884ms       64.28016ms            5.26%
  - initCCL             |    0.02914942ms        3.614528ms           0.296%
  - mergeCCL            |    0.08448568ms        10.47622ms           0.857%
  - compressCCL         |    0.02113755ms        2.621056ms           0.214%
bboxes                  |     0.7925288ms        98.27357ms            8.04%
Mem. Management         |      1.211777ms        150.2604ms            12.3%

Start to finish: 1.22s
```

GPU V1.1

Changes:
- Background is processed only once
- Memory is allocated and free only once (based on frame dimensions)

Effect:
Only 2 Memcpy per frame and overhaul performance improvement
⅓ of memory management time
~+30 FPS

```
--------------------------------------------------------------------------------
OPENCV Bench (v1):
FPS: 413.4

STEP                |      FRAME_AVG          TOTAL        EXEC_TIME

grayscale           |      243.7581µs        30226µs          9.87%
blur                |       340.121µs        42175µs          13.8%
diff                |      17.23387µs         2137µs         0.698%
threshold           |      11.10484µs         1377µs          0.45%
morph               |      1140.677µs       141444µs          46.2%
connectedComps      |      200.5968µs        24874µs          8.12%
bboxes              |      115.3548µs        14304µs          4.67%

Start to finish: 306ms
--------------------------------------------------------------------------------
GPU Bench (v1.1):
blockDim: 32x32
gridDim: 19x11
FPS: 135.5

STEP                |      FRAME_AVG          TOTAL        EXEC_TIME

grayscaleGPU        |   0.03420232ms       4.241088ms          0.461%
blurGPU             |     1.527637ms        189.427ms           20.6%
  - getGaussianMatrix | 3.019355e-05ms       0.003744ms       0.000407%
diffGPU             |   0.02343923ms       2.906464ms          0.316%
thresholdGPU        |   0.02263226ms        2.8064ms           0.305%
morph               |     4.235504ms       525.2025ms            57%
  - dilateGPU       |     2.109465ms       261.5737ms          28.4%
  - erodeGPU        |     2.111233ms       261.7929ms          28.4%
  - getCircleKernel | 9.290323e-06ms       0.001152ms       0.000125%
connectedComps      |    0.4274418ms        53.00278ms          5.76%
  - initCCL         |   0.02929265ms       3.632288ms          0.394%
  - mergeCCL        |   0.04559897ms       5.654272ms          0.614%
  - compressCCL     |   0.01731381ms       2.146912ms          0.233%
bboxes              |       0.6278ms        77.8472ms          8.45%
Mem. Management     |    0.4903283ms        60.8007ms           6.6%

Start to finish: 0.921s
```

GPU V1.2

Changes:
- Shared memory tiling for blur and morph
- Input/Output tile version differs drastically
- Constant memory for kernel masks

Effect:
First sm version double exec time
Second version halved it
~+60 FPS

```
--------------------------------------------------------------------------------
OPENCV Bench (v1):
FPS: 310.2

STEP                    |      FRAME_AVG          TOTAL          EXEC_TIME

grayscale               |     365.1452µs         45278µs             11.2%
blur                    |     469.0323µs         58160µs             14.3%
diff                    |     31.21774µs          3871µs            0.954%
threshold               |     18.41935µs          2284µs            0.563%
morph                   |     1209.121µs        149931µs               37%
connectedComps          |     310.3306µs         38481µs             9.48%
bboxes                  |     215.6129µs         26736µs             6.59%

Start to finish: 406ms
--------------------------------------------------------------------------------
GPU Bench (v1.2):
blockDim: 32x32
gridDim: 19x11
FPS: 192.8

STEP                    |      FRAME_AVG          TOTAL          EXEC_TIME

grayscaleGPU            |     46.05626µs        5.710976ms            0.88%
blurGPU                 |     911.9634µs        113.0835ms            17.4%
  - getGaussianMatrix   |   0.05909677µs        0.007328ms          0.00113%
diffGPU                 |     28.38865µs        3.520192ms            0.543%
thresholdGPU            |       25.088µs        3.110912ms            0.479%
morph                   |     2275.744µs        282.1922ms            43.5%
  - dilateGPU           |     1122.045µs        139.1336ms            21.4%
  - erodeGPU            |     1126.151µs        139.6428ms            21.5%
  - getCircleKernel     |   0.01832258µs        0.002272ms          0.00035%
connectedComps          |     596.8764µs        74.01267ms            11.4%
  - initCCL             |     34.73961µs        4.307712ms            0.664%
  - mergeCCL            |     60.76542µs        7.534912ms            1.16%
  - compressCCL         |     19.18348µs        2.378752ms            0.367%
bboxes                  |     457.4072µs         56.7185ms            8.74%
Mem. Management         |     711.3737µs        88.21034ms            13.6%

Start to finish: 0.649s
```

```
OPENCV Bench (v1):
FPS: 456.5

STEP                    |      FRAME_AVG           TOTAL        EXEC_TIME

grayscale               |     189.7661µs          23531µs           8.49%
blur                    |     330.1694µs          40941µs           14.8%
diff                    |     15.49194µs           1921µs          0.693%
threshold               |     11.78226µs           1461µs          0.527%
morph                   |      1093.46µs         135589µs           48.9%
connectedComps          |     170.6613µs          21162µs           7.64%
bboxes                  |     106.5645µs          13214µs           4.77%

Start to finish: 277ms
-------------------------------------------------------------------------------
GPU Bench (v1.2):
blockDim: 32x32
gridDim: 19x11
FPS: 94.29

STEP                    |      FRAME_AVG           TOTAL        EXEC_TIME

grayscaleGPU            |     32.65832µs        4.049632ms          0.306%
blurGPU                 |     2781.915µs         344.9574ms           26.1%
  - getGaussianMatrix   |   0.09806452µs          0.01216ms        0.00092%
diffGPU                 |     23.98658µs         2.974336ms          0.225%
thresholdGPU            |     22.22555µs         2.755968ms          0.208%
morph                   |     6747.217µs         836.6549ms          63.3%
  - dilateGPU           |     3367.165µs         417.5284ms          31.6%
  - erodeGPU            |     3367.104µs         417.5209ms          31.6%
  - getCircleKernel     | 0.008774194µs         0.001088ms       8.23e-05%
connectedComps          |     383.2534µs         47.52342ms          3.59%
  - initCCL             |     28.09652µs         3.483968ms          0.264%
  - mergeCCL            |     14.47871µs          1.79536ms          0.136%
  - compressCCL         |     14.51613µs             1.8ms          0.136%
bboxes                  |     156.5272µs         19.40938ms          1.47%
Mem. Management         |     480.1894µs         59.54349ms           4.5%

Start to finish: 1.32s
```

First version of the shared memory load

GPU V1.3

Changes:

- Simple search for optimal block dimensions
- From 4 to 32-wide-blocks, 16 are the most efficient

Effect:

First sm version double exec time

Second version halved it

~+30 FPS

```
--------------------------------------------------------------------------------
OPENCV Bench (v1):
FPS: 408.8

STEP                |      FRAME_AVG          TOTAL        EXEC_TIME

grayscale           |     222.8145µs        27629µs            8.93%
blur                |     363.2742µs        45046µs            14.6%
diff                |     16.83871µs         2088µs           0.675%
threshold           |     13.43548µs         1666µs           0.538%
morph               |     1157.871µs       143576µs            46.4%
connectedComps      |     186.9032µs        23176µs            7.49%
bboxes              |     111.2016µs        13789µs            4.46%

Start to finish: 309ms
--------------------------------------------------------------------------------
GPU Bench (v1.2):
blockDim: 16x16
gridDim: 38x21
FPS: 224

STEP                |      FRAME_AVG          TOTAL        EXEC_TIME

grayscaleGPU        |     32.48748µs       4.028448ms           0.721%
blurGPU             |     907.5257µs       112.5332ms            20.1%
  - getGaussianMatrix | 0.03225807µs          0.004ms         0.000716%
diffGPU             |      24.2031µs       3.001184ms           0.537%
thresholdGPU        |       24.48µs        3.03552ms           0.543%
morph               |     2246.633µs       278.5825ms            49.8%
  - dilateGPU       |      1115.96µs        138.379ms            24.8%
  - erodeGPU        |     1118.135µs       138.6487ms            24.8%
  - getCircleKernel | 0.01522581µs         0.001888ms         0.000338%
connectedComps      |      413.247µs       51.24262ms            9.17%
  - initCCL         |     23.17419µs        2.8736ms           0.514%
  - mergeCCL        |     46.39794µs       5.753344ms            1.03%
  - compressCCL     |     20.41806µs        2.53184ms           0.453%
bboxes              |     362.9897µs       45.01072ms            8.05%
Mem. Management     |     478.1806µs        59.2944ms            10.6%

Start to finish: 0.559s
```

Memory Usage

All memory allocation at launch:

- 2x 3-channel image [width * height * sizeof(uchar3)] ⇒ original background and input, one is free before processing
- 3x single-channel frame [width * height * sizeof(uchar3)] ⇒ current bgd, frame and swap frame

Memory bandwidth at launch

- 2x Memcpy to constant [ksize * ksize * sizeof(float|uchar)] ⇒ blur and morph kernels
- 1x Memcpy for background processing (3 frames == 1 image)

Memory bandwidth for each frame

- 1x Memcpy for original frame (~3 frames)
- 1x Memcpy to receive labelled frame [width * height * sizeof(**int**)]

nvprof - 16x16 blocks

Kernel	achieved_occupancy	sm_efficiency	branch_efficiency	shared_efficiency	shared_utilization
	%	%	%		
grayscaleGPU	0.812064	84.27	100.00	0.00	Idle (0)
blurTiledConstantGPU	0.738575	98.99	100.00	24.96	Mid (4)
grayscaleGPU	0.810232	85.65	100.00	0.00	Idle (0)
blurTiledConstantGPU	0.739220	99.06	100.00	24.96	Mid (4)
diffGPU	0.785808	82.60	100.00	0.00	Idle (0)
thresholdGPU	0.806869	66.39	100.00	0.00	Idle (0)
dilateTiledConstantG	0.717766	99.15	100.00	24.93	Low (2)
erodeTiledConstantGP	0.717814	99.13	100.00	24.93	Low (2)
initCCL	0.810692	83.75	100.00	0.00	Idle (0)
mergeCCL	0.773982	79.46	100.00	0.00	Idle (0)
compressCCL	0.743263	75.47	100.00	0.00	Idle (0)

What is well-optimized

- **Memory coalescing**, accesses are always coalesced
- **Memory communication** unless bboxes generation can be done on the GPU
- **Kernel launch overhead**, all kernels are executed sequentially, and there is not much to be done here

Future versions

- Shared memory prefetching for the blur and morph kernels
- Loop unrolling, specifically for memory loading
- Bounding boxes on GPU

Thank you for your attention!